

UVA CS 4774: Machine Learning

Lecture 12: Neural Network (NN) and More: BackProp


Dr. Yanjun Qi

University of Virginia

Department of Computer Science

Last: Logistic Regression **Classifier**

$$P(y = 1|x) = \frac{e^{\beta_0 + \beta^T x}}{1 + e^{\beta_0 + \beta^T x}} \quad p(\text{Head})$$




Last: Logistic Regression Classifier


- View I: logit(y) as linear of Xs
- View II: model Y as Bernoulli with $p(y=1|x)$ as $p(\text{Head})$
- View III: S" shape function compress to $[0,1]$
- View IV: models a linear classification boundary!
- View V: Two stages: summation + sigmoid

$$\ln \left[\frac{P(y|x)}{1-P(y|x)} \right] = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

$$P(y=1|x) = \frac{e^{\beta_0 + \beta^T x}}{1 + e^{\beta_0 + \beta^T x}}$$




Today: Basic Neural Network Models



Roadmap: DNN Basics

- Basics of Neural Network (NN)



- single neuron, e.g. logistic regression unit
- multilayer perceptron (MLP)
- various loss function
 - E.g., when for multi-class classification, softmax layer
- training NN with backprop algorithm
 - A few advanced tricks

ReWrite Logistic Regression as two stages:


First:

Summing $z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$

Second:

Sigmoid
Squashing $\hat{y} = P(y=1|x) = \frac{e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p}}{1 + e^{\beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p}} = \frac{e^z}{1 + e^z}$

One “Neuron”: Expanded Logistic Regression






$$z = w^T \cdot x + b$$

$$y = \text{sigmoid}(z) = \frac{e^z}{1 + e^z}$$

E.g., Many Possible Nonlinearity Functions

(aka transfer or activation functions)





Name	Plot	Equation	Derivative (w.r.t x)
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
Rectifier (ReLU) ^[9]		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$

Sigmoid is any S-shaped function

usually works best in practice


Activation functions

In machine learning, “sigmoid” usually refers to this logistic function, so the terms are often used interchangeably.

Name and Graph	Function	Derivative
sigmoid(x) 	$h(x) = \frac{1}{1 + \exp(-x)}$	$h'(x) = h(x)[1 - h(x)]$
tanh(x) 	$h(x) = \frac{\exp(x) - \exp(-x)}{\exp(x) + \exp(-x)}$	$h'(x) = 1 - h(x)^2$
softplus(x) 	$h(x) = \log(1 + \exp(x))$	$h'(x) = \frac{1}{1 + \exp(-x)}$
rectify(x) 	$h(x) = \max(0, x)$	$h'(x) = \begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$


History → Perceptron: 1-Neuron Unit with Step

- First proposed by Rosenblatt (1958)
- A simple neuron that is used to classify its input into one of two categories.
- A perceptron uses a **step function**




Sigmoid


$$\sigma(x) = 1/(1 + e^{-x})$$




tanh tanh(x)



ReLU max(0,x)




Leaky ReLU
 $\max(0.1x, x)$




Maxout $\max(w_1^T x + b_1, w_2^T x + b_2)$

ELU
$$f(x) = \begin{cases} x & \text{if } x > 0 \\ \alpha (\exp(x) - 1) & \text{if } x \leq 0 \end{cases}$$



Bias Term?




$$z = w^T \cdot x + b$$

$$y = \text{sigmoid}(z) = \frac{e^z}{1 + e^z}$$

Without Bias Term


$$z = \alpha x$$

$\text{sig}(z)$ ↑



→ x


With Bias Term



Roadmap: DNN Basics


- Basics of Neural Network (NN)
 - single neuron, e.g. logistic regression unit
 - ➔ • multilayer perceptron (MLP)
 - various loss function
 - E.g., when for multi-class classification, softmax layer
 - training NN with backprop algorithm

Neuron Representation




The linear transformation and nonlinearity together is typically considered a single neuron

Multi-Layer Perceptron (MLP)- (Feed-Forward NN)



Multi-Layer Perceptron (MLP)- (Feed-Forward NN)




$$\begin{aligned}
 z_1 &= W_1^T x \\
 h_1 &= \text{sigmoid}(z_1) \\
 z_2 &= W_2^T h_1 \\
 h_2 &= \text{sigmoid}(z_2) \\
 z_3 &= W_3^T h_2 \\
 \hat{y} &= \text{sigmoid}(z_3)
 \end{aligned}$$

hidden layer 1 output


hidden layer 2 output

$$x \xrightarrow{w_1} z_1 \rightarrow h_1 \xrightarrow{w_2} z_2 \rightarrow h_2 \xrightarrow{w_3} \hat{y}$$

“Deep” Neural Networks (i.e. many hidden layers)



Revolution of Depth




View IV: Logistic Regression models a linear classification boundary!

$$y \in \{0,1\}$$

$$\ln \left[\frac{P(y|x)}{1 - P(y|x)} \right] = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

$p(y=1|x)$
 $= p(y=0|x)$
 $= 0.5$
Boundary
no decision





Decision Boundary → equals to zero



$$\ln \left[\frac{P(y=1|x)}{P(y=0|x)} \right] = \ln \left[\frac{P(y=1|x)}{1 - P(y=1|x)} \right] = \alpha + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

$= w^T \vec{x} + b = 0$

Activation function



Activation function











[3]

$$f(x) = \frac{1}{1 + e^{-x}}$$


$x \rightarrow \Sigma \rightarrow \text{sig} \rightarrow f(x)$

$\text{Sigmoid}(z) \rightarrow \hat{y} = f(x)$





Today: Basic Neural Network Models



Thank You



UVA CS 4774: Machine Learning

Lecture 12: Neural Network (NN) and More: BackProp

Dr. Yanjun Qi

University of Virginia

Department of Computer Science


Module II

Roadmap: DNN Basics

- Basics of Neural Network (NN)
 - single neuron, e.g. logistic regression unit
 - multilayer perceptron (MLP)
 -  • various loss function
 - E.g., when for multi-class classification, softmax layer
 - training NN with backprop algorithm


E.g., SSE loss on Multi-Layer Perceptron (MLP) for Regression

Example: 2 Hidden Layer MLP network with 2 output units:



$$E_w(\hat{\mathbf{y}}, \mathbf{y}) = \text{SSE} = (y_1 - \hat{y}_1)^2 + (y_2 - \hat{y}_2)^2$$

e.g., Cross-Entropy loss for Multi-Layer Perceptron (MLP) for Binary Classification




$$E_x(\theta) = Loss_x(\theta) = -\log \Pr(Y = y | X = x) \\ = -\{y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)\}$$

?

Cross-entropy loss function, OR named as "deviance", OR negative log-likelihood

e.g., Cross-Entropy loss for Multi-Layer Perceptron (MLP) for Binary Classification



LIKELIHOOD:

Basic Bernoulli

$$L(p) = \prod_{i=1}^n p^{z_i} (1-p)^{1-z_i}$$

↑
function of $p = \text{Pr}(\text{head})$

$$\begin{aligned} \log(L(p)) &= \log \left[\prod_{i=1}^n p^{z_i} (1-p)^{1-z_i} \right] \\ &= \sum_{i=1}^n (z_i \log p + (1-z_i) \log(1-p)) \end{aligned}$$

Logistic / Bernoulli

$$\begin{aligned} L(\beta) &= \prod_{i=1}^n p(y_i=1|x_i) (1-p(y_i=1|x_i))^{1-y_i} \\ &= \prod_{i=1}^n \hat{y}_i^{y_i} (1-\hat{y}_i)^{1-y_i} \end{aligned}$$

Log likelihood

$$\begin{aligned} \ell(\beta) &= \sum_{i=1}^n \left(y_i \log \hat{y}_i \right. \\ &\quad \left. + (1-y_i) \log(1-\hat{y}_i) \right) \end{aligned}$$

Binary Classification → Multi-Class Classification

models the **target binary random** variable with Bernoulli whose parameter $p=p(y=1|x)$ predefined as function on x




$1-p(y=1|x)$

$$P(y=1|x) = \frac{e^z}{1 + e^z}$$

$p(c|x)$

Multinoulli Distribution, e.g.,



k

$$\left\{ \begin{array}{l} p(y=1|x) \\ p(y=2|x) \\ p(y=3|x) \\ \vdots \\ p(y=6|x) \end{array} \right.$$

Multi-class target variable representation

- Multi-class output variable →

An indicator basis vector representation

- If output variable **G** has **K** classes, there will be **K indicator variable** y_i

$K=4$
Total Class

y_k

$k=4$

one-hot

g	y₁	y₂	y₃	y₄
3	0	0	1	0
1	1	0	0	0
2	0	1	0	0
4	0	0	0	1
1	1	0	0	0

$$P(y = c_k | x)$$

$G: \{1, 2, 3, 4\} \rightarrow k=4$

$C: \{\text{poli}, \text{fin}, \text{Health}, \text{Enter}, \text{tech}\} \rightarrow k=5$

Review: Multi-class variable representation

Class y_k

g	y_1	y_2	y_3	y_4
3	0	0	1	0
1	1	0	0	0
2	0	1	0	0
4	0	0	0	1
1	1	0	0	0

N

- Multi-class output variable →
An indicator basis vector representation
 - If output variable G has K classes, there will be K indicator variable y_i
- How to classify to multi-class ?
 - First: learn K different regression (✗)
 - Then: Softmax using all K outputs as input

Review: Multi-class variable representation

Class

g	y₁	y₂	y₃	y₄
3	0	0	1	0
1	1	0	0	0
2	0	1	0	0
4	0	0	0	1
1	1	0	0	0

N

- Multi-class output variable →
An indicator basis vector representation
 - If output variable **G** has **K** classes, there will be **K** indicator variable y_i

• How to classify to multi-class ?

- First: learn K different regression
- Then: Softmax using all K outputs as input

- Then: $\hat{G}(x) = \underset{k \in g}{\operatorname{argmax}} \hat{f}_k(x)$

Identify the largest component of $\hat{f}(x)$
And Classify according to MAP Rule

Review: Multi-class variable representation

Class y_k

g	y_1	y_2	y_3	y_4
3	0	0	1	0
1	1	0	0	0
2	0	1	0	0
4	0	0	0	1
1	1	0	0	0

N

- Multi-class output variable $\rightarrow$
An indicator basis vector representation
 - If output variable G has K classes, there will be K indicator variable y_i


- How to classify to multi-class ?
 - First: learn K different regression
 - Then: Softmax using all K outputs as input
 - Then:

$$\hat{G}(x) = \underset{k \in g}{\operatorname{argmax}} \underbrace{\hat{f}_k(x)}_{p_k(x)}$$

Identify the largest component of $\hat{f}(x)$
And Classify according to MAP Rule

$$p(y=k|x)$$


discriminative
classifier



When for multi-class classification
(last output layer: softmax layer)

last layer is softmax output layer →
a Multinoulli logistic regression unit

Strategy : Use “softmax” layer function for multi-class classification




$$Pr(G = k | X = x) = Pr(Y_k = 1 | X = x)$$

$$y_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$$

“Softmax” function: Normalizing function which converts each class output to a probability.

$$\frac{\partial y_i}{\partial z_i} = y_i (1 - y_i)$$




When for multi-class classification
(last output layer: softmax layer)

last layer is softmax output layer →
a Multinoulli logistic regression unit


y_{i1}	y_{i2}	y_{i3}
0	1	0
1	0	0
0	0	1


class 2
class 1
class 3

multiclass
logistic regression



$$\hat{y}_i = \frac{e^{z_i}}{\sum_j e^{z_j}} = P(y_i = 1 \mid x)$$






When for multi-class classification
(last output layer: softmax layer)

last layer is softmax output layer →
a Multinoulli logistic regression unit

$$E_W(\hat{y}, y) = \text{cross-E} = - \sum_{j=1 \dots K} y_j \ln \hat{y}_j$$

MLE / the negative log probability of the right
answer / Cross entropy loss function :




When for multi-class classification
(last output layer: softmax layer)

last layer is softmax output layer $\rightarrow$
a Multinoulli logistic regression unit

$$E_W(\hat{y}, y) = \text{cross-E} = - \sum_{j=1 \dots K} y_j \ln \hat{y}_j$$

MLE / the negative log probability of the right
answer / Cross entropy loss function :



When for multi-class classification
(last output layer: softmax layer)

last layer is softmax output layer →
a Multinoulli logistic regression unit

$$E_W(\hat{y}, y) = \text{cross-E} = - \sum_{j=1 \dots K} y_j \ln \hat{y}_j$$


MLE / the negative log probability of the right
answer / Cross entropy loss function :

$$\frac{\partial E}{\partial z_i} = \sum_{j=1 \dots K} \frac{\partial E}{\partial y_j} \frac{\partial y_j}{\partial z_i} = \hat{y}_i - y_{true,i}$$

Error calculated from
predicted Output vs. true

Summary Recap: Multi-Class Classification Loss

Cross Entropy Loss



$$\hat{y}_i = \frac{e^{z_i}}{\sum_j e^{z_j}} = P(y_i = 1 \mid x)$$

“Softmax” function.
Normalizing function which
converts each class output to
a probability.

$$\begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \quad \begin{pmatrix} 0.1 \\ 0.7 \\ 0.2 \end{pmatrix}$$

y $\hat{y}$

$$E = \text{loss} = - \sum_{j=1 \dots K} y_j \log \hat{y}_j$$

“0” for all except true class


Logistic: a special case of softmax for two classes

$\{H, T\}$

$$y_1 = \frac{e^{z_1}}{e^{z_1} + e^{z_0}} = \frac{1}{1 + e^{-(z_1 - z_0)}}$$

- So the logistic binary case is just a special case that avoids using redundant parameters:
 - Adding the same constant to both z_1 and z_0 has no effect.
 - The over-parameterization of the softmax is because the probabilities must add to 1.

Today: Basic Neural Network Models



Thank You



UVA CS 4774: Machine Learning

Lecture 12: Neural Network (NN) and More: BackProp

Dr. Yanjun Qi

University of Virginia


Department of Computer Science

Module III


Roadmap: DNN Basics

- Basics of Neural Network (NN)
 - single neuron, e.g. logistic regression unit
 - multilayer perceptron (MLP)
 - various loss function
 - E.g., when for multi-class classification, softmax layer
-  • training NN with backprop algorithm
 - A few advanced tricks


e.g., “Block View” of Logistic Regression




e.g., “Block View” of Logistic Regression



e.g., “Block View” of multi-class NN



$$x \xrightarrow{W_1} h_1 \xrightarrow{W_2} h_2 \xrightarrow{W_3} \hat{y} \rightarrow E(\hat{y}, y)$$



Review: Stochastic GD →

- For LR: linear regression, We have the following descent rule:

$$\theta_j^{t+1} = \theta_j^t - \alpha \frac{\partial}{\partial \theta_j} \underline{J(\theta)} \Big|_t$$

$E(\hat{y}, y)$



- → For neural network, we have the delta rule

$$\Delta \mathbf{w} = -\eta \frac{\frac{\partial E}{\partial \mathbf{w}^t}}$$

$$\mathbf{w} = \{w^1, w^2, \dots, w^T\}$$

$$\underline{W^{t+1} = W^t - \eta \frac{\partial E}{\partial W^t} = W^t + \Delta w}$$

Backpropagation

- 1. Initialize network with random weights
- 2. For training examples:
 - **Forward:** Feed feed inputs to network layer by layer, and calculate output of each layer (from input layer to until the final layer (error function))
 - **Backward:** For all layers (starting with the output layer, back to input layer):
 - Propagate local gradients layer by layer from final layer, until back to input layer to calculate each layer's gradient $\frac{\partial E}{\partial W_l^t}$  Need to calculate these!
 - Adapt weights in current layer
$$W_l^{t+1} = W_l^t - \eta \frac{\partial E}{\partial W_l^t}$$

Training Neural Networks by Backpropagation - to jointly optimize all parameters


How do we learn the optimal weights W_L for our task??

- Stochastic Gradient descent:

$$W_L^{t+1} = W_L^t - \eta \frac{\partial E_x}{\partial W_L^t}$$

But how do we get gradients of lower layers?



- Backpropagation!
 - Repeated application of chain rule of calculus
 - Locally minimize the objective
 - Requires all “blocks” of the network to be differentiable




Layers of Differentiable Parameterized Functions (with nonlinearities)

Forward: Feed inputs to network layer by layer, and calculate output of each layer
(from input layer to until the final layer (error function))

Need to calculate these!



Layers of Differentiable Parameterized Functions (with nonlinearities)



Backward: For all layers (starting with the output layer, back to input layer):

Propagate local gradients layer by layer from final layer, until back to input layer to calculate each layer's gradient

$$\underbrace{\frac{\partial E}{\partial W_l^t}} \longrightarrow \text{Need to calculate these!}$$

Adapt weights in current layer

$$\underbrace{W_l^{t+1} = W_l^t - \eta \frac{\partial E}{\partial W_l^t}}$$

Training Neural Networks by Backpropagation - to jointly optimize all parameters


How do we learn the optimal weights W_L for our task??

- Stochastic Gradient descent:

$$W_L^{t+1} = W_L^t - \eta \frac{\partial E}{\partial W_L^t}$$

But how do we get gradients of lower layers?


- Backpropagation!
 - Repeated application of chain rule of calculus
 - Locally minimize the objective
 - Requires all “blocks” of the network to be differentiable




$$E(x, y) = f_4(y, f_3(f_2(f_1(x)))$$

$W_3 \quad W_2 \quad W_1$


“Local-ness” of Backpropagation




“Local-ness” of Backpropagation




Example: Sigmoid Block



Example: Softmax Block (right before loss layer)



Backpropagation Example



f_1	$z_1 = x_1 w_1 + x_2 w_3 + b_1$ $z_2 = x_1 w_2 + x_2 w_4 + b_2$
f_2	$h_1 = \frac{\exp(z_1)}{1 + \exp(z_1)}$ $h_2 = \frac{\exp(z_2)}{1 + \exp(z_2)}$
f_3	<u>$\hat{y} = h_1 w_5 + h_2 w_6 + b_3$</u>
f_4	$E = (y - \hat{y})^2$

argmin_w { $f_4 (f_3 (f_2 (f_1 ()))$) }

$$\frac{\partial E}{\partial w_5} = \frac{\partial E}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w_5}$$

$$= \underline{-2(y - \hat{y})} \underline{h_1}$$

$$\frac{\partial E}{\partial w_5} =$$

$$\frac{\partial E}{\partial w_1} =$$

Extra

$\operatorname{argmin}_w \{f_4(f_3(f_2(f_1(\cdot))))\}$

f_1	$z_1 = x_1 w_1 + x_2 w_3 + b_1$ $z_2 = x_1 w_2 + x_2 w_4 + b_2$
f_2	$h_1 = \frac{\exp(z_1)}{1 + \exp(z_1)}$ $h_2 = \frac{\exp(z_2)}{1 + \exp(z_2)}$
f_3	$\hat{y} = h_1 w_5 + h_2 w_6 + b_3$
f_4	$E = (y - \hat{y})^2$

Input

Output

Local Gradients= ∂ Output / ∂ Input


.

$\text{argmin}_w \{f_4(f_3(f_2(f_1(\cdot))))\}$


f_1	$z_1 = x_1 w_1 + x_2 w_3 + b_1$ $z_2 = x_1 w_2 + x_2 w_4 + b_2$
f_2	$h_1 = \frac{\exp(z_1)}{1 + \exp(z_1)}$ $h_2 = \frac{\exp(z_2)}{1 + \exp(z_2)}$
f_3	$\hat{y} = h_1 w_5 + h_2 w_6 + b_3$
f_4	$E = (y - \hat{y})^2$

Input	Output	Local Gradients = $\partial \text{Output} / \partial \text{Input}$
$x_1, x_2, w_1, \dots$	z_1, z_2	$\frac{\partial z_1}{\partial x_1} = w_1$
z_1, z_2	h_1, h_2	$\frac{\partial h_1}{\partial z_1} = h_1(1-h_1)$
w_5, h_1, h_2	$\hat{y}$	$\partial \hat{y} / \partial h_1 = w_5$
$\hat{y}$	loss E	$\partial E / \partial \hat{y} = -2(y - \hat{y})$


$$\begin{aligned}
 \frac{\partial E}{\partial w_1} &= \frac{\partial E}{\partial w_1} = \frac{\partial f_4}{\partial f_3} \frac{\partial f_3}{\partial f_2} \frac{\partial f_2}{\partial f_1} \frac{\partial f_1}{\partial w_1} \\
 &= -2(y - \hat{y}) \frac{\partial (h_1 w_5 + h_2 w_6 + b_3)}{\partial w_1} \\
 &= -2(y - \hat{y}) \left(w_5 \frac{\partial h_1}{\partial w_1} + w_6 \frac{\partial h_2}{\partial w_1} \right) \\
 &= -2(y - \hat{y}) w_5 \frac{\partial h_1}{\partial w_1} = -2(y - \hat{y}) w_5 \frac{\partial h_1}{\partial z_1} \frac{\partial z_1}{\partial w_1} \\
 &= \underbrace{-2(y - \hat{y})}_{f_4 \text{ local}} \underbrace{w_5}_{f_3 \text{ local}} \underbrace{h_1(1-h_1)}_{f_2 \text{ local}} \underbrace{x_1}_{\frac{\partial f_1}{\partial w_1}}
 \end{aligned}$$



gradients to the
next previous layer





BackProp in Practice: Mini-batch SGD




1. Initialize weights
2. For each batch of input samples Sx :
 - a. Run the network “Forward” on S to compute outputs and loss
 - b. Run the network “Backward” using outputs and loss to compute gradients
 - c. Update weights using SGD (or a similar method)
2. Repeat step 2 until loss convergence

Monitor and visualize the loss curve



Gradient Magnitudes:



Gradients too big $\rightarrow$ divergence


Gradients too small $\rightarrow$ slow convergence

Divergence is much worse!

Many great tools, e.g., Adam

<https://arxiv.org/abs/1609.04747>

Monitor and visualize the train / validation loss / accuracy: Bias Variance Tradeoff



big gap = overfitting

=> increase regularization strength?

no gap, e.g. underfitting / both bad

=> increase model capacity?

Other things to plot and check:

- Per-layer activations:
 - Magnitude, center (mean or median), breadth (sdev or quartiles)
 - Spatial/feature-rank variations
- Gradients
 - Magnitude, center (mean or median), breadth (sdev or quartiles)
 - Spatial/feature-rank variations
- Learning trajectories
 - Plot parameter values in a low-dimensional space

Hyperparameters to play with:

- network architecture
- learning rate, decay schedule, update type
- regularization (L2/Dropout strength)

How to become a great neural networks practitioner

→ Craft? / Talent? / Experience?

Your Friend: loss function





Weight Initialization


```
input layer had mean 0.001800 and std 1.001311
hidden layer 1 had mean 0.001198 and std 0.627953
hidden layer 2 had mean -0.000175 and std 0.486051
hidden layer 3 had mean 0.000055 and std 0.407723
hidden layer 4 had mean -0.000306 and std 0.357108
hidden layer 5 had mean 0.000142 and std 0.320917
hidden layer 6 had mean -0.000389 and std 0.292116
hidden layer 7 had mean -0.000228 and std 0.273387
hidden layer 8 had mean -0.000291 and std 0.254935
hidden layer 9 had mean 0.000361 and std 0.239266
hidden layer 10 had mean 0.000139 and std 0.228008
```

```
W = np.random.randn(fan_in, fan_out) / np.sqrt(fan_in) # layer initialization
```

“Xavier initialization”
[Glorot et al., 2010]



Reasonable initialization.
(Mathematical derivation
assumes linear activations)



Batch Normalization: implicit regularization

[Ioffe and Szegedy, 2015]

Normalize:

$$\hat{x}^{(k)} = \frac{x^{(k)} - \mathbb{E}[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$


And then allow the network to squash the range if it wants to:

$$y^{(k)} = \gamma^{(k)} \hat{x}^{(k)} + \beta^{(k)}$$

- Improves gradient flow through the network
- Allows higher learning rates
- Reduces the strong dependence on initialization
- Acts as a form of regularization in a funny way, and slightly reduces the need for dropout, maybe


Standardizing the activations of the prior layer means that assumptions the subsequent layer makes about the spread and distribution of inputs during the weight update will not change, at least not dramatically. This has the effect of stabilizing and speeding-up the training process of deep neural networks.

Data Preprocessing




Regularization by Dropout

“randomly set some neurons to zero in the forward pass”



(a) Standard Neural Net




(b) After applying dropout.

[Srivastava et al., 2014]

Dropout is training a large ensemble of models (that share parameters).
Each binary mask is one model

Dropout At test time....




Ideally:

want to integrate out all the noise

Monte Carlo approximation:

do many forward passes with different dropout masks, average all predictions

Today: Basics of Neural Network Models



Thank You



References

- ❑ Dr. Yann Lecun's deep learning tutorials
- ❑ Dr. Li Deng's ICML 2014 Deep Learning Tutorial
- ❑ Dr. Kai Yu's deep learning tutorial
- ❑ Dr. Rob Fergus' deep learning tutorial
- ❑ Prof. Nando de Freitas' slides
- ❑ Olivier Grisel's talk at Paris Data Geeks / Open World Forum
- ❑ Hastie, Trevor, et al. *The elements of statistical learning*. Vol. 2. No. 1. New York: Springer, 2009.
- ❑ Dr. Hung-yi Lee's CNN slides

UVA CS 4774: Machine Learning

Lecture 12: Neural Network (NN) and More: BackProp


Dr. Yanjun Qi

University of Virginia

Department of Computer Science


Module IV

“Local-ness” of Backpropagation




Backpropagation

(binary classification example)



Backpropagation

(binary classification example)



$$E = \text{loss} = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

**Gradient
Descent to
Minimize loss:**


$$\mathbf{w}_2(t + 1) = \mathbf{w}_2(t) - \eta \frac{\partial E}{\partial \mathbf{w}_2(t)}$$

$$W_1(t + 1) = W_1(t) - \eta \frac{\partial E}{\partial W_1(t)}$$

Need to find
these!

Backpropagation

(binary classification example)



$$E = f_4 = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = f_3 = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = f_3 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = f_2 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = f_1 = \mathbf{W}_1^T \mathbf{x}$$

$$E = f_4(f_3(f_2(f_1(x))))$$

Backpropagation

(binary classification example)



$$E = f_4 = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = f_3 = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = f_3 = \mathbf{w}_2^T \mathbf{h}_1$$

$$\mathbf{h}_1 = f_2 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = f_1 = \mathbf{W}_1^T \mathbf{x}$$


$$\frac{\partial E}{\partial \mathbf{w}_2} = ??$$

$$\frac{\partial E}{\partial \mathbf{W}_1} = ??$$

$$E = f_4(f_3(f_2(f_1(x))))$$

Backpropagation

(binary classification example)



$$E = f_4 = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = f_3 = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = f_3 = \mathbf{w}_2^T \mathbf{h}_1$$

$$\mathbf{h}_1 = f_2 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = f_1 = \mathbf{W}_1^T \mathbf{x}$$


$$\frac{\partial E}{\partial \mathbf{w}_2} = ??$$

$$\frac{\partial E}{\partial \mathbf{W}_1} = ??$$

$$E = f_4(f_3(f_2(f_1(x)))) \longrightarrow \text{Exploit the chain rule!}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = w_2^T h_1$$


$$h_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$z_1 = W_1^T x$$

$$\frac{\partial E}{\partial w_2} =$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\frac{\partial E}{\partial \mathbf{w}_2} = \overbrace{\frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{w}_2}}^{\text{chain rule}}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{w}_2} &= \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{w}_2} \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot \end{aligned}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{w}_2} &= \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{w}_2} \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot \left(\frac{e^{z_2}}{1 + e^{z_2}} \left(1 - \frac{e^{z_2}}{1 + e^{z_2}} \right) \right) \end{aligned}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{w}_2} &= \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{w}_2} \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot \left(\frac{e^{z_2}}{1 + e^{z_2}} \left(1 - \frac{e^{z_2}}{1 + e^{z_2}} \right) \right) \cdot (\mathbf{h}_1) \end{aligned}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = W_1^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{w}_2} &= \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{w}_2} \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot \left(\frac{e^{z_2}}{1 + e^{z_2}} \left(1 - \frac{e^{z_2}}{1 + e^{z_2}} \right) \right) \cdot (\mathbf{h}_1) \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot (\hat{y}(1 - \hat{y})) \cdot (\mathbf{h}_1) \end{aligned}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$z_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\frac{\partial E}{\partial \mathbf{W}_1} =$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\frac{\partial E}{\partial \mathbf{W}_1} = \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{h}_1} \cdot \frac{\partial \mathbf{h}_1}{\partial z_1} \cdot \frac{\partial z_1}{\partial \mathbf{W}_1}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$


$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

$$\begin{aligned} \frac{\partial E}{\partial \mathbf{W}_1} &= \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{h}_1} \cdot \frac{\partial \mathbf{h}_1}{\partial \mathbf{z}_1} \cdot \frac{\partial \mathbf{z}_1}{\partial \mathbf{W}_1} \\ &= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot (\hat{y}(1 - \hat{y})) \cdot (\mathbf{w}) \cdot (\mathbf{h}_1(1 - \mathbf{h}_1)) \cdot (\mathbf{x}) \end{aligned}$$

Backpropagation

(binary classification example)



$$E = -y \ln(\hat{y}) - (1 - y) \ln(1 - \hat{y})$$

$$\hat{y} = \frac{e^{z_2}}{1 + e^{z_2}}$$

$$z_2 = \mathbf{w}_2^T \mathbf{h}_1$$

$$\mathbf{h}_1 = \frac{e^{z_1}}{1 + e^{z_1}}$$

$$\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$$

already computed!

$$\frac{\partial E}{\partial \mathbf{W}_1} = \frac{\partial E}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial z_2} \cdot \frac{\partial z_2}{\partial \mathbf{h}_1} \cdot \frac{\partial \mathbf{h}_1}{\partial \mathbf{z}_1} \cdot \frac{\partial \mathbf{z}_1}{\partial \mathbf{W}_1}$$

$$= \left(\frac{\hat{y} - y}{\hat{y}(1 - \hat{y})} \right) \cdot (\hat{y}(1 - \hat{y})) \cdot (\mathbf{w}) \cdot (\mathbf{h}_1(1 - \mathbf{h}_1)) \cdot (\mathbf{x})$$